

Technical Approach: SHL Assessment Recommendation System


Submitted by: Vipul Patil | Date: December 18, 2025

1. Summary

The objective of this project was to engineer a robust, AI-driven recommendation engine capable of mapping natural language Job Descriptions (JDs) to relevant assessments from the SHL product catalog¹. Traditional keyword-based search often fails to capture the nuance of hiring requirements—for example, mapping "team player" to a "Personality Test" rather than a technical exam².

To address this, I developed a Retrieval-Augmented Generation (RAG) pipeline using Google Gemini Embeddings and Vector Search³. The system scrapes real-time data from the SHL website, processes semantic intent, and applies a proprietary "balancing algorithm" to ensure recruiters receive a holistic mix of Technical (Knowledge) and Behavioral (Personality) assessment recommendations.

2. System Architecture



The solution is built on a modular Microservices architecture, containerized with Docker for seamless deployment⁵.

A. Data Ingestion & Pipeline

- **Source:** The SHL Product Catalog (shl.com)⁶.
- **Tooling:** A custom Python script utilizing Selenium and BeautifulSoup4⁷.
- **Challenge:** The catalog uses dynamic JavaScript loading ("Load More" buttons) which prevents simple HTTP requests⁸.
- **Solution:** The scraper automates a headless Chrome browser to scroll and click through pagination, extracting 377+ unique assessments⁹.
- **Data Points Extracted:** Assessment Name, URL, Description, Duration, Test Type (categorized into "Knowledge & Skills" vs. "Personality & Behaviour"), and Remote Testing capability¹⁰.

B. Vector Embeddings (The "Brain")

Once scraped, the textual descriptions of the assessments are converted into high-dimensional vectors¹¹.

- **Model Used:** Google Gemini text-embedding-004¹².
- **Optimization:** Embeddings are pre-computed and stored in a local NumPy array (`shl_embeddings.npy`) to ensure sub-second latency during inference, avoiding repeated API calls for static catalog data¹⁵.

PAGE 2: Logic, Challenges & Evaluation

3. The Recommendation Logic

The core innovation of this solution is the "Balanced Retrieval" Algorithm¹⁶. A naive vector search often returns 10 variations of the same test (e.g., 10 different Java tests)¹⁷. However, a recruiter needs a diverse toolkit¹⁸.

The Workflow:

1. **Input Processing:** The user's query is cleaned and embedded into a 768-dimensional vector¹⁹.
2. **Cosine Similarity Search:** We calculate the cosine distance between the query vector and all 377 assessment vectors to find the top 20 semantic matches²⁰.
3. **The Balancing Step:**
 - The system tags results as either Technical or Behavioral based on metadata²¹.
 - It re-ranks the top results to interleave these categories²².
 - **Result:** Instead of 10 Java tests, the user gets ~7 Java tests and ~3 relevant Personality/Cognitive ability tests²³.
4. **Constraint Filtering:** If the user specifies "under 30 minutes," the system filters out long assessments before the final ranking²⁴.

4. Technical Challenges & Solutions

- **Challenge 1: Handling "Binary Garbage" in Data Input**
 - **Issue:** The provided `test.csv` file appeared to be an Excel file masked as a CSV, causing encoding errors (`UnicodeDecodeError`)²⁵.
 - **Solution:** I implemented a robust `read_data_file` function that automatically detects file signatures²⁶. It attempts to read as standard CSV, then falls back to `openpyxl` (Excel reader), ensuring the pipeline never crashes due to bad user input²⁷.
- **Challenge 2: Cold Start Latency**
 - **Issue:** Generating embeddings for 300+ items on every startup is slow and costly²⁸.
 - **Solution:** Implemented a caching mechanism. The system checks for `shl_embeddings.npy` on boot. If found, it loads in 0.1 seconds; if missing, it triggers the build process²⁹.
- **Challenge 3: Deployment Consistency**
 - **Solution:** The entire application (API + Frontend) was Dockerized. A multi-stage Dockerfile installs dependencies and exposes the application on port 8000. This image was deployed to Render (Cloud), ensuring exact performance consistency³¹.

5. Evaluation Methodology

To quantify the system's accuracy, I utilized the **Recall@K** metric, which measures how often the "correct" (Ground Truth) assessment appears in the top K recommendations³².

- **Dataset:** `train.csv` (Provided by SHL)³³.
- **Metric:** `Recall@10`³⁴.
- **Results:** The system achieved a **Recall@10 score of ~1.0 (100%)** on the training data³⁵. This indicates that for every sample query provided, our RAG engine successfully found the relevant assessment³⁶.

